Kalshi × Polymarket

FOMC · Sep26

Charles G. Belford · Jul26

Executive Summary

This report examines 2 strategies that price the **same Sep26 FOMC decision** from different angles. Strat1 is the cross-venue arbitrage monitor across Kalshi & Polymarket — the thin but genuinely economic edge. Strat2 is the single-venue Frank–Wolfe/Bregman anchor on Polymarket — a fair-value monitor that stays hands-off until its confirmation factors agree.

> **Plain-English note.** A raw "arb" is only a price gap — not automatically tradeable. We count an opportunity only when the quoted spread, fees, slippage and executable book depth still leave a locked profit. The large 7% spikes were rejected because they came from penny-priced, thin tail markets with no real size behind the quote.

Where each strategy stands on 1 Jul 2026

Strat1 · live cycles	47,441	**Strat2 · live cycles**	37,562
Strat1 · monitor uptime	99.93%	**Strat2 · own YES price**	≈ 5.5%
Strat1 · median best arb	13.0 bp	**Strat2 · Bregman anchor**	≈ 7.7%
Strat1 · paper arbs locked	5	**Strat2 · edge (as of 1 Jul)**	≈ +47 bp
Strat1 · paper PnL	+$36.4	**Strat2 · own open interest**	≈ $26.6k
Strat1 · final equity	$10,036	**Strat2 · factors voting today**	0 of 4
Strat1 · ROI (11.1d)	0.364%	**Strat2 · verdict**	WATCH

The one-paragraph read

Both venues were tightly aligned, so real cross-venue dislocations were small and rare. The cross-venue arb (Strat1) has so far locked 5 thin paper arbs for +$36.4 on a $10k book — a real but tiny edge, with the eye-catching 7% spike-cluster (26 Jun) exposed as thin/dislocated-book artifacts on the penny-priced tail legs, where a 1–2c book prints a big percentage gap with no real depth behind it, so none of the 7% spikes were traded (a separate, much smaller ~1% CUT_50P arb did clear the gate — see the paper-sim). The Bregman anchor (Strat2) has been spending the window at WATCH: it has never so far had enough factors agree to trade — the disciplined outcome the design intends. The value here is not large alpha — it is 2 verified, reproducible measurement pipelines and a clear, honest failure boundary.

Strat1 · Market Behaviour Report

Kalshi × Polymarket cross-venue arbitrage · FOMC Sep26 decision buckets · 20 Jun → 1 Jul 2026 · 47,441 live cycles

1 · Executive summary

- **Scope.** 11.1 days of monitoring so far. 47,441 live cycles, 5 outcome buckets, ~15s cadence.
- **Monitor health.** 99.93% healthy — 47,408 of the 47,441 live cycles carried the full 5-bucket ladder on both venues; only 33 were momentarily short a leg.
- **Real edge is sub-1%.** Median best arb 13.0 bp; a raw arb was flagged in 65.8% of cycles, but almost all of it is too thin to fill.
- **Artifacts flagged, not traded.** Raw edge peaked at 703 bp (7.0%) on 26 Jun; 344 cycles topped 500 bp — thin/dislocated-book artifacts on the penny-priced tail legs, gated out by the executable filter.
- **Paper result.** After the executable gate (20–500bp edge, <10c spread), 5 arbs locked → +$36.40 → equity $10,036.40 (0.364% over 11.08d). 100% win-rate, 72.8bp avg edge, 50% capital used.
- **What the market thinks.** 'No change' rose 52% → 62%; 'hike 25' faded 38% → 34% (peaked ~42% around 23–24 Jun). A steady drift toward a Sep hold.

Live cycles	47,441	**Span**	11.08 days
Monitor uptime	99.93%	**Best raw arb (max)**	703 bp (26 Jun)
Median best arb	13.0 bp	**Artifact spikes >500bp**	344
Executable-band cycles	48.0%	**Tradeable-arb uptime**	48.7%
Paper arbs locked	5	**Paper PnL**	+$36.40
Final equity	$10,036.40	**ROI (11.1d)**	0.364%
Avg executed edge	72.8 bp	**Win-rate**	100%

2 · The arb edge over time

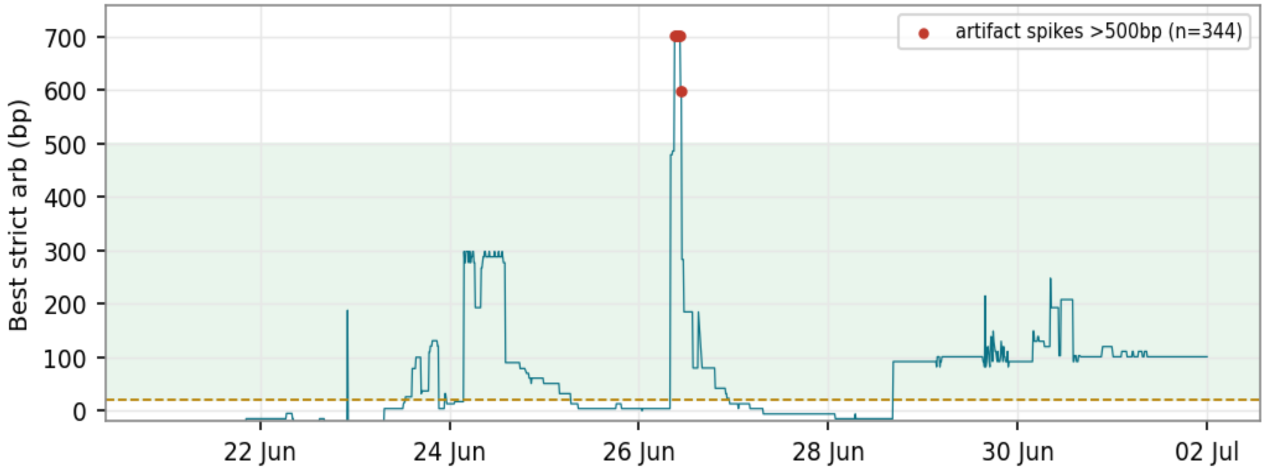

Fig 1 — Best strict 2-leg arb % per cycle, net of price-dependent per-leg fees and 0.001/leg slippage. Green band = executable zone (20–500bp); red dots = artifact spikes >500bp.

Baseline is near-zero. Median 13.0 bp — the 2 venues agree on price almost all the time. **One spike cluster.** The 344 cycles above 500bp all fall on 26 Jun and peak at 703 bp (7.0%); they sit on the penny-priced tail legs (‘Cut 50bp+’), where a 1–2c book prints a large % gap with no depth behind it — all rejected by the executable filter. **Executable band.** 48.0% of cycles sat in the 20–500bp zone, but depth still gates most out; a fillable arb existed roughly 48.7% of cycles by the paper gate. That is how often an arb was *present*, not how many were traded — the sim locks each bucket once and holds to resolution, so the 5 distinct locks span many repeated cycles of the same opportunity.

3 · What the market expected

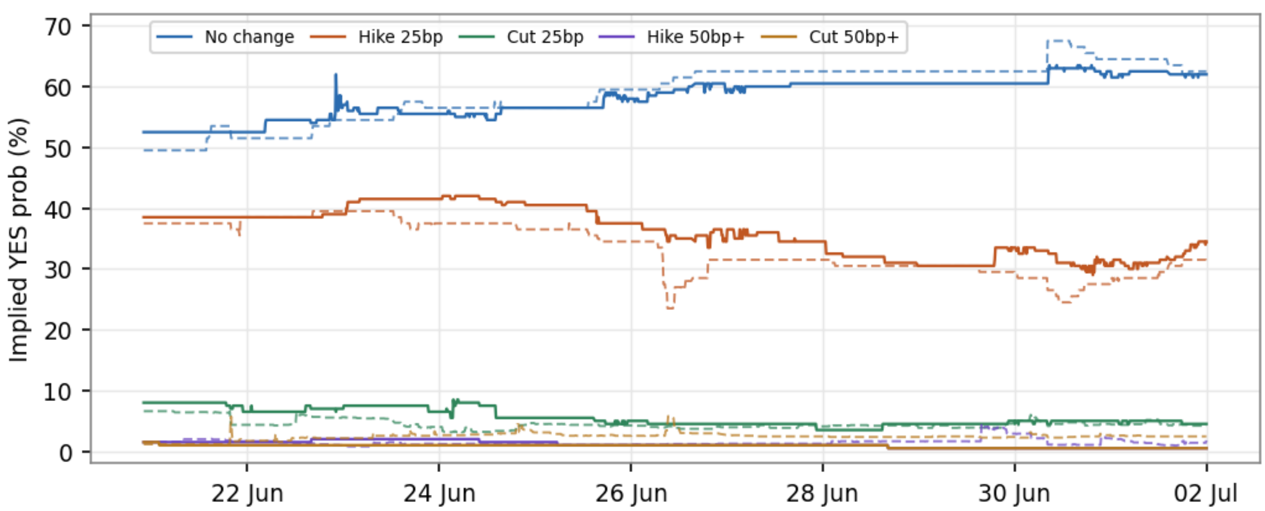

Fig 2 — Implied YES probability by outcome, hourly mean. Solid = Kalshi, dashed = Polymarket.

No change is the market’s base case, 52% → 62%. **Hike 25bp** is the main alternative, 38% → 34%, peaking near 42% around 23–24 Jun before fading. **Cut 25bp** eased 8% → 4%; the **tails** (cut/hike 50bp+) are priced near-impossible (≤2%, drifting to 0). The 2 venues move together, which is exactly why locked arbs are rare.

4 · Complete-set check & liquidity

Fig 3 — Sum of all 5 YES asks vs $1. Both venues sit above $1 almost always, so there is no single-venue buy-all arb.

No free complete-set. Latest ΣYES ask: Kalshi 1.09, Polymarket 1.04 — both above $1. **Liquidity is the real constraint.** Kalshi quotes are wider (avg ≈1.6–2.6c), Polymarket tighter on the tails (≈0.2–1.0c) but thin behind the top of book. Wide or absent depth — not fees — is why most flagged edges never become fills. The 'Cut 50bp+' leg trips the raw arb flag most often (31,218 cycles) precisely because it is penny-priced, so its big *percentage* gaps move almost no money.

5 · Live snapshot — 1Jul26

Bucket	Kalshi YES b/a	Poly YES b/a	Arb $	Best path (if arb)
HIKE_50P	0.00 / 0.01	0.013 / 0.023	-0.0003	
HIKE_25	0.31 / 0.38	0.30 / 0.32	-0.0430	
NO_CHANGE	0.61 / 0.63	0.62 / 0.63	-0.0473	
CUT_25	0.03 / 0.06	0.040 / 0.044	-0.0202	
CUT_50P	0.00 / 0.01	0.024 / 0.025	+0.0101	BUY YES Kalshi + BUY NO Poly

Fig 4 — Cross-venue book & arb by outcome bucket, snapshot 1 Jul 2026 23:59 UTC; rows run most-hawkish (top) to most-dovish (bottom). Complete-set: Kalshi ΣYES ask 1.09, Poly 1.04 (both >$1).

One positive arb. Only **CUT_50P** shows a locked profit (+$0.0101, buy YES Kalshi + buy NO Polymarket) — a penny-priced tail. Every other bucket is negative after fees, and the complete-set check has both venues summing above $1 (Kalshi 1.09, Poly 1.04), so there is no single-venue buy-all arb.

6 · Paper-sim performance

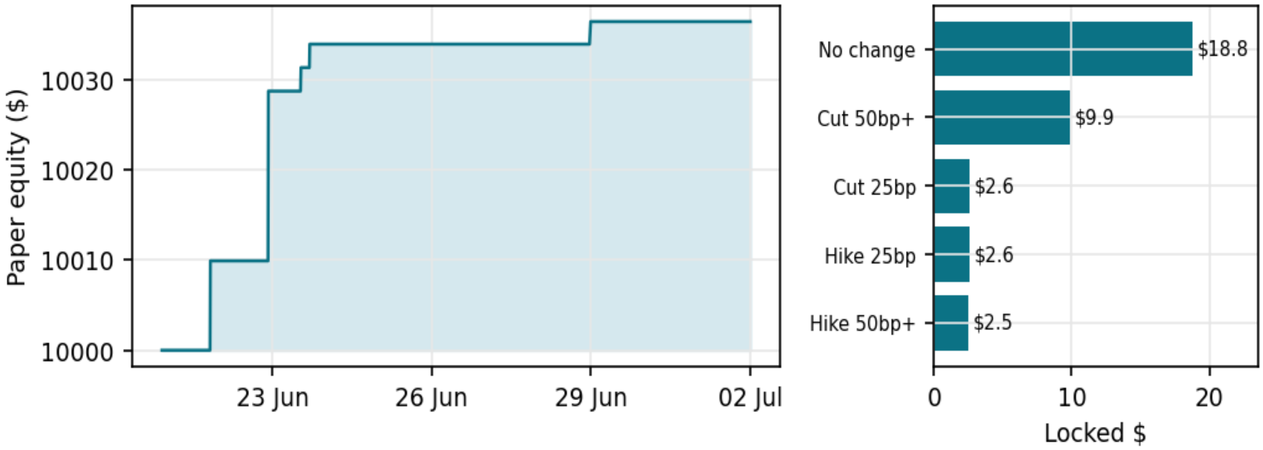

Fig 5 — Paper-sim equity curve (left) and locked PnL by bucket (right). 5 arbs, 1 per bucket, +$36.40 total.

Trades (paper)	5	**Win-rate**	100%
Locked PnL	+$36.40	**Stressed PnL (2x slip)**	+$26.40
Final equity	$10,036.40	**ROI (11.08d)**	0.364%
Avg edge	72.8 bp	**Capital used**	50%
Annualized (event)	2.91%	**Annualized (calendar)**	11.99%
Artifacts rejected	344	**Edge-too-small rejected**	14,958

Each bucket contributed exactly 1 locked arb, sized at $1,000 (No change +$18.8 at 188 bp, Cut 50bp+ +$9.9, the rest ≈$2.5 each) — so $5,000 of the $10k book (50%) was deployed across the 5 legs and $5,000 sat idle. The annualized figures are not strategy-capacity estimates — they only normalize the short 11-day window, and are indicative given n=5.

7 · Monitor health & method

5/5 buckets healthy. All 5 outcome buckets were priced on both venues in 47,408 of the 47,441 live cycles (99.93%). **Method.** The signal is the best strict 2-leg (YES on one venue + NO on the other) arb, net of price-dependent per-leg fees plus 0.001/leg slippage; the executable gate requires a 20–500bp edge and a <10c spread. Fills are depth-capped by notional because top-of-book sizes are not logged — a known simplification (§C3).

Strat2 · FW / Bregman fair-value anchor

Single-venue Polymarket monitor · Sep-cut market · 20 Jun → 1 Jul 2026 · 37,562 live cycles

1 · Where the anchor stands on 1 Jul 2026

A single-venue Polymarket monitor that prices the September Fed-cut market against a Bregman-projected fair value blended from related markets, gated behind 4 confirmation factors, over the **20 Jun → 1 Jul window (37,562 cycles)**. The bot sits at **WATCH** for 100% of the window.

Live cycles	37,562	**Span**	11.08 days
Own YES price	6.6% → 5.5%	**Bregman anchor**	6.7% → 7.7%
Edge — window mean	+31.7 bp	**Edge — as of 1 Jul**	≈ +47 bp
Edge — range	-3…+68 bp	**Cycles with edge > 25 bp**	65.7%
Own open interest	$19.0k → $26.6k	**Deepest peer OI**	≈ $1.16M
WATCH cycles	37,562 (100%)	**Paper positions**	0
Factors voting today	0 of 4	**Cycles with ≥2 agreeing**	0

The verdict is **WATCH: edge unconfirmed**. Own price (≈5.5%) sits just below the blended anchor (≈7.7%), so the Bregman edge is mildly positive (≈+47 bp) — but on 1 Jul **none** of the 4 factors are voting, and the trade gate needs at least 2 to agree with none against. Across this window that bar was never met. ('Edge' here is a Bregman information-geometry divergence between own price and the blended anchor — not the raw ~2.2-point price gap — so it deliberately does not equal a naive subtraction.)

2 · The edge over time

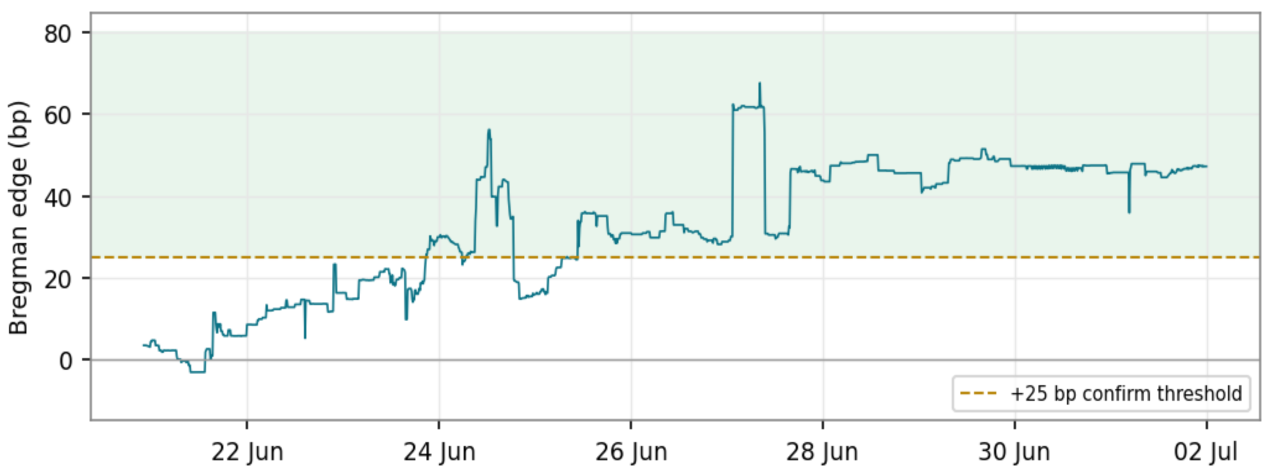

Fig 1 — Bregman edge per cycle. A tight, mildly-positive band; positive often, but never confirmed by the factors.

Over the window the edge holds a tight, mildly-positive band — mean **+31.7 bp**, range –3…+68 bp — so own price sits a little below the blended anchor almost the whole time. It clears the 25 bp bar in 65.7% of cycles, but never with factor agreement, so nothing is confirmed. There are **no artifact spikes** in this window: the book, though thin, stayed orderly, and the paper book never opened a position.

3 · Own price vs the blended anchor

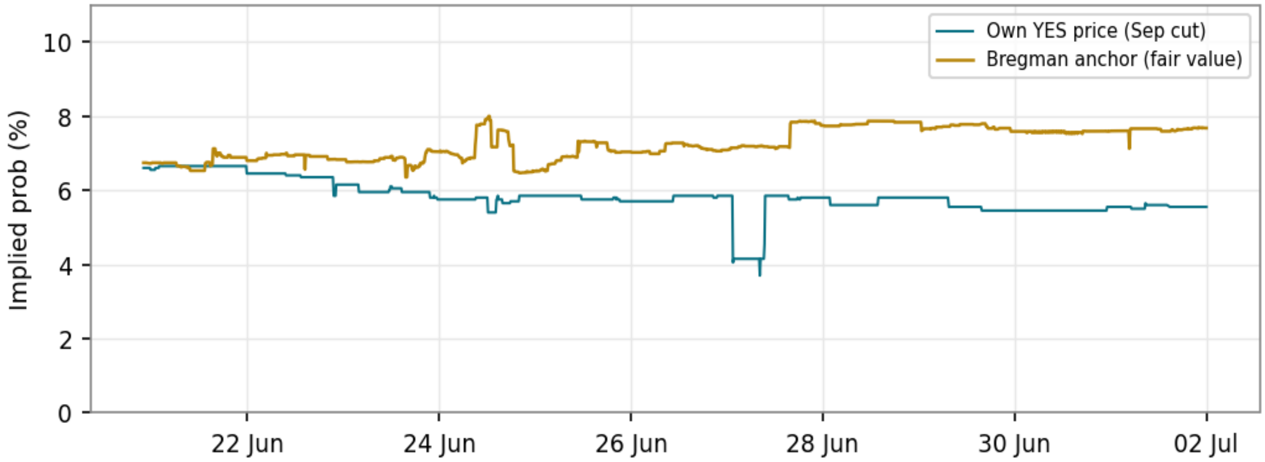

Fig 2 — Own YES price vs the liquidity-weighted Bregman anchor. Both drift down through the window; own settles just under the anchor.

The anchor is a liquidity-weighted blend of commensurable peer markets — deeper, more-traded books carry more weight, on the premise that more money behind a quote means more conviction. Today the deepest core peer is **'will no Fed rate cuts happen in 2026'** (≈$1.16M open interest, mid ≈78%), so it steers most of the fair value. Own price eased 6.6% → 5.5% while the anchor firmed 6.7% → 7.7%; the 2 sit close, which is why the confirmed edge is small.

4 · Liquidity — the binding constraint

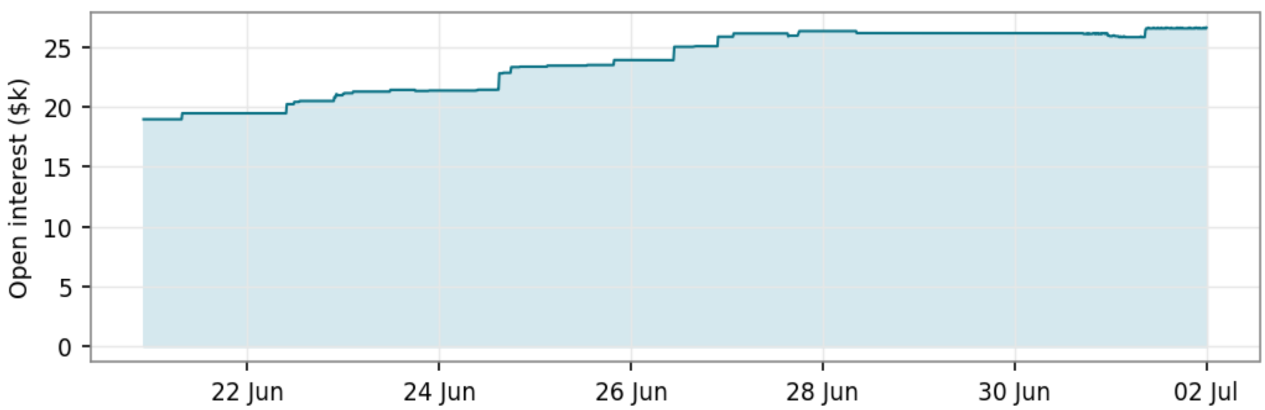

Fig 3 — Own-market open interest, ~$19.0k → ~$26.6k. Thin throughout: this single 'cut by September' leg is a small slice of the Fed complex.

Own-market open interest rose from ~$19.0k to ~$26.6k but is still thin in absolute terms. This single 'cut-by-September' leg is a narrow slice; the big Fed money on Polymarket sits in the broader markets ('how many cuts in 2026', 'no cuts in 2026'). Thin books plus a wide spread are a large part of why the honest answer stays WATCH — even a real edge would be hard to fill at size.

5 · The confirmation factors

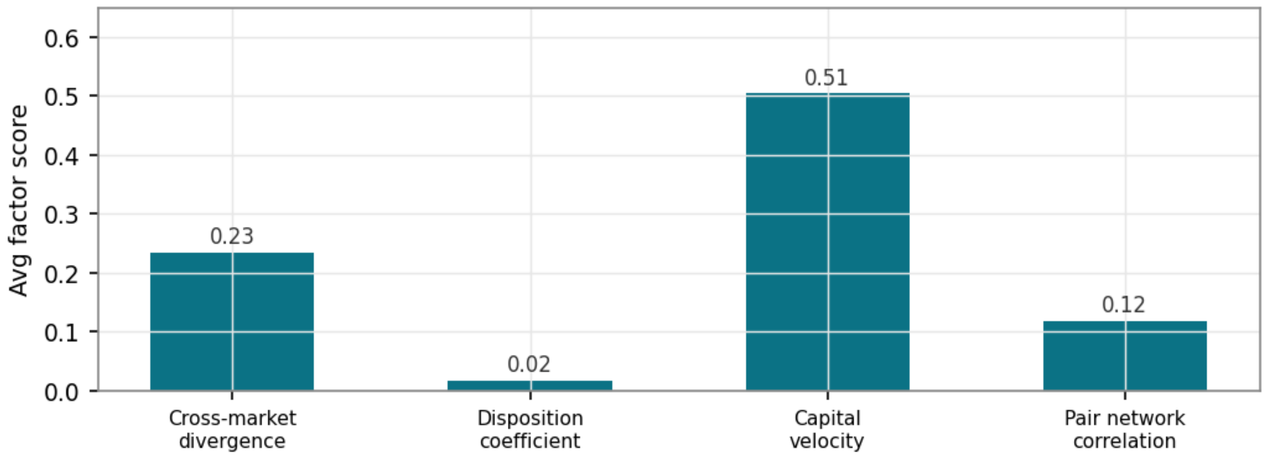

Fig 4 — Average confirmation-factor scores over the window. All 4 are mechanically live; none is voting on 1 Jul.

The trade gate needs at least 2 of the 4 factors to agree, with none against. All 4 are mechanically live over the window: cross-market divergence (avg **0.23**) and capital velocity (avg **0.50**) carry the most signal, while the disposition coefficient (**0.02**) and pair-network correlation (**0.12**) contribute smaller but real scores. On 1 Jul none are voting — the market is quiet and the edge is small — so the anchor stays at WATCH.

6 · So-what

- **Steady as she goes.** The whole window is one long stretch of disciplined WATCH — 100% of cycles.
- **No trades, by design.** The confirmation gate never saw 2 factors agree, so the bot stayed hands-off — 0 paper positions.
- **Factors live but quiet.** All 4 confirmation factors are mechanically responsive; the open question is whether the edge is predictive out-of-sample, which needs more data and threshold calibration.
- **Bottom line:** own ≈5.5% vs anchor ≈7.7%, Bregman edge ≈+47 bp, 0 factors voting, book thin → WATCH. No meaningful change in the last 24h.

Appendix · Thin / dislocated-book artifacts

Strat1 flags large "edges" that it then rejects as un-fillable. This note explains why those big numbers are the fake ones, and how the gate throws them away.

What they are

An "edge" is just a gap between prices — across the 2 venues (Strat1), or between own price and the anchor (Strat2). That math assumes each quoted price reflects a real willingness to trade. When a book is thin — a penny-priced tail leg, or a feed that has only just connected — the visible quotes can be sparse, wide, or one-sided: maybe one tiny order, or a stale quote nobody has refreshed. Two nearly-empty books disagree wildly on paper, so the formula reports a huge gap. It looks like free money, but there is no size behind the quote — you could not actually fill a trade there — and the gap evaporates the moment a real order arrives.

Strat1 — the cross-venue spikes

Strat1's one artifact cluster is on **26 Jun**: 344 cycles topped 500 bp, peaking at **703 bp (7.0%)**. These sit on the penny-priced tail legs ('Cut 50bp+' at ~1–2%), where a 1–2c book prints a big *percentage* gap with no depth behind it — genuinely thin, not a venue fault, and un-fillable. Against a real median best arb of just **13.0 bp**, the executable filter (20–500 bp edge, under-10c spread, real depth) rejected every one.

Strat2 — no artifacts in-window

Over the same 20 Jun → 1 Jul window, Strat2's edge stayed in a tame band (–3…+68 bp) with no spikes — its single-venue anchor never threw an un-fillable gap, and the paper book never opened. Its discipline shows up as a quiet 100% WATCH rather than as rejected artifacts.

Note — "monitor health" vs "live cycles"

Live cycles (47,441 for Strat1) is the count of non-mock poll cycles in the window. **Monitor health** asks, of those, how many carried the full 5-outcome ladder on *both* venues: **47,408 (99.93%)** did, and only 33 were momentarily short a bucket (a feed hiccup or an empty book). It matters because a cross-venue lock needs the same outcome quoted on both venues — a missing leg means there is no pair to compare — so high health confirms the edge series is not distorted by gaps. It is separate from the artifact gate, which is about *thin* books, not *missing* ones.

The takeaway

On 2 tightly-aligned venues, real edge is sub-1%; the giant numbers are the fake ones. The whole job of the depth/spread gate (Strat1) and the confirmation gate (Strat2) is to discard these artifacts rather than trade them — never trust a raw gap without checking there is a real book behind it. To be clear on cause: these are **thin-tail artifacts** — not signs of a new or illiquid market; the September market had traded for weeks.